

The Evolution of Command-and-Control (C2) Traffic and NDR Blind Spots

Executive Research Article & Comprehensive Technical Report

Document ID	SOC-NDR-MASTER-001-EN
Author	Oğulcan (ByGhost) Erarslan
Title	Senior Solutions Architect Cyber Security Specialist
Organization	ByGhost Security
Website	https://byghost.tr/
GitHub	https://github.com/ByGh00st
Contact	linkedin.com/in/byghost-tr
Target Audience	SOC Leadership, Blue Team Leads, Security Architects
Scope	Executive Research Article + 12-Section Comprehensive Technical Report
Methodology	Stochastic timing analysis, NDR architecture, endpoint telemetry correlation
Date	October 04, 2026
Volume	9,699 words

© 2026 Oğulcan (ByGhost) Erarslan. This technical report is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0) license. Commercial integration into security products (EDR/NDR/XDR) or managed service tiers without explicit written authorization is strictly prohibited. Refer to Appendix E for full legal notice and intellectual property specifications.

Table of Contents

PART I: EXECUTIVE RESEARCH ARTICLE

1. The Jitter Fallacy: Why ~20% Random Jitter Fails as an Evasion Primitive
2. The Evolution: Self-Similar Traffic and the Collapse of Flow-Level Discrimination
3. The Paradigm Shift: Contextual and Causal Detection Engineering

PART II: COMPREHENSIVE TECHNICAL REPORT

Executive Summary

1. Scope, Assumptions, and Threat Model
 2. Mathematical Foundations of Timing Analysis
 3. Statistical Fingerprint of Enterprise Traffic
 4. Scope Boundary and Non-Operational Model Context
 5. The Scope of the Detection Ceiling
 6. Tactical Catalog (T1 – T20)
 7. Detection Engineering Deliverables
 8. False Positive Management and Quantitative Metrics
 9. Maturity Model and Implementation Roadmap
 10. Defensive Debt and Continuous Calibration
 11. Conclusion
 12. Appendices (A – E)
- Appendix E — Legal Notice & Intellectual Property Armor

PART I: EXECUTIVE RESEARCH ARTICLE

The Evolution of Command-and-Control Traffic: Shaped Timing Dynamics and NDR Blind Spots

> LEGAL AND INTELLECTUAL PROPERTY NOTICE (CC BY-NC-ND 4.0):

> "This research publication and technical specification is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0). It may not be used or reproduced without explicit author attribution; nor may it be integrated, directly or indirectly, into commercial cybersecurity products (EDR/NDR/XDR), managed services, or proprietary fee-based offerings without prior written authorization and an express commercial licensing agreement from the author. No derivative works or modifications are permitted. Document Identifier: SOC-NDR-MASTER-001-EN / Author: Oğulcan Erarslan (ByGhost)."

1. The Jitter Fallacy: Why ~20% Random Jitter Fails as an Evasion Primitive

A widespread assumption across defensive and adversarial engineering asserts that adding random jitter—on the order of 20%—to beacon intervals reliably evades network detection systems. This assumption is structurally misleading: it conflates two fundamentally different stochastic point-process models and misdirects detection architecture design.

The dichotomy of two distinct jitter models. In the first model, stochastic noise is injected directly into absolute event timestamps ($t_n = nT + \epsilon_n$): because consecutive inter-arrival times (IAT) are derived from adjacent event times ($IAT_n = t_n - t_{n-1} = T + \epsilon_n - \epsilon_{n-1}$), the resulting IAT sequence inherently exhibits shared noise terms and statistical dependency. Under the formal conditions that (i) the noise sequence is identically and independently distributed (i.i.d.), (ii) possesses finite variance, and (iii) the event indices preserve physical chronological ordering, this dependency assumes an invariant structure: the lag-1 autocorrelation reaches -0.5 . In scenarios where packet loss, transmission jitter, or transport-layer reordering disrupts chronological sequencing, this invariance cannot be guaranteed. In the second model, stochastic perturbation is added directly to the interval duration itself ($IAT_n = T + U_n$): intervals are mutually independent, and absolute event times cumulatively form a renewal process ($t_n = \sum IAT_i$). These two models yield fundamentally different marginal IAT distributions, dependency structures, and spectral profiles; detection logic must explicitly state which stochastic model is assumed.

To which object does the spectrum belong? The discrete Fourier transform (DFT) of a constant, scalar IAT sequence produces exclusively a DC component ($f = 0$); discrete harmonic spectral lines originate from the continuous-time spectrum of the periodic point event train (the Dirac comb), not from the scalar IAT sequence itself. A detector that applies spectral analysis exclusively to scalar IAT arrays will structurally fail to detect even a rigid, fixed-period beacon. The impact of jitter on this harmonic line structure is strictly model-dependent: under the event-time jitter model, harmonic line amplitudes are scaled by the characteristic function (Fourier transform) of the jitter distribution, and harmonics coinciding with the zeros of this envelope are attenuated or suppressed entirely; under the interval-jitter model, a continuous interval distribution completely eliminates discrete line components for $f \neq 0$, transitioning the spectrum into a purely continuous density. In the case of uniform event-time jitter, the spectral envelope takes the form of $\text{sinc}(2\pi T f)$ where $\text{sinc}(x) \equiv \sin(x)/x$; this expression is a direct consequence of characteristic function theory and contains no arbitrary system parameters. As observation duration and event sample sizes increase, the statistical power of spectral estimation scales according to the chosen windowing and tapering methodologies.

Observations of bounded interval support $[T_{\min}, T_{\max}]$ are likewise context-sensitive: legitimate enterprise applications, depending on protocol class and workload profile, exhibit substantially wider and highly diverse IAT supports; lower bounds are routinely imposed by protocol handshakes, transport buffering, thread schedulers, and timestamp clock resolution. Therefore, bounded support is not conclusive proof of malicious behavior, but merely

a candidate feature. Similarly, timing, spectral, and distributional summaries can be fed into machine learning-based anomaly detection models; however, the empirical performance of such models depends entirely on feature representation, training set diversity, and baseline stationarity. Asserting that any specific architecture will inherently succeed without empirical validation across real-world workloads is technically ungrounded.

Section Conclusion: In the event-time jitter model, the harmonic line structure of a periodic event train degrades but does not vanish unconditionally; in the interval-jitter model, a continuous interval density eliminates discrete spectral lines for $f \neq 0$. Consequently, the evasion efficacy of jitter is strictly governed by the underlying stochastic model. Defensive detection frameworks must formulate hypotheses with explicitly defined stochastic assumptions.

2. The Evolution: Self-Similar Traffic and the Collapse of Flow-Level Discrimination

Self-similarity and long-range dependence have been observed in many measured aggregated Ethernet/WAN and enterprise traffic traces, with $H > 0.5$ reported in those LRD regimes. Under temporal aggregation, memoryless Poisson processes flatten asymptotically into white noise, whereas LRD processes retain memory across multiple temporal scales. This empirical observation cannot be attributed to a single cause: while heavy-tailed source distributions (such as Pareto-distributed file transfers) provide a primary mechanism, multi-scale multiplexing, application-layer request clustering, and network queueing artifacts also contribute. It is an empirical finding across measured traces, not a universal property of all enterprise traffic.

When an adversary's C2 timing profile adapts to an enterprise baseline, cross-class discrimination within classical flow-level feature families—marginal IAT distributions, power spectral density, autocorrelation decay, scaling exponents, and diurnal rhythms—can diminish drastically. This is not an information-theoretic claim of absolute absence: higher-order statistical dependencies, multifractal structure, protocol state consistency, cross-flow multi-host correlation, and host endpoint context retain discriminative capacity. The attenuation of flow-level separation does not mean that detection has become impossible; it dictates that detection efforts must migrate to orthogonal feature families.

3. The Paradigm Shift: Contextual and Causal Detection Engineering

Human-interactive workloads and automated processes exhibit distinct temporal and contextual variability patterns; however, this distinction is not absolute and requires workload-specific baselines. When discrimination collapses within flow-level statistics, defensive value concentrates across the following orthogonal candidate features:

- **Scaling Stability:** The moving-window variance of scaling exponents (H) and tail index estimates evaluated against segment baselines; anomalous parameter stability serves as a candidate anomaly signal requiring multi-plane correlation.
- **Tail Structure:** Systematic deviations in the empirical upper tail relative to the segment baseline, diagnosed through tail-index estimators (e.g., Hill estimator) and complementary cumulative distribution function (CCDF) plots.
- **Timer Quantization & Continuous Sampling:** Clustering artifacts induced by operating system thread schedulers, runtime sleep APIs, power management policies, and hardware timestamp clock sources; these represent candidate features evaluated against host, build, and segment baselines. Conversely, evaluating whether continuous float-domain sampling ($T = -\ln(U) / \lambda$) has been introduced to break discrete timer buckets.
- **Fleet-Wide Similarity:** Unusually uniform timing or scaling parameters observed across disparate endpoints within an enterprise fleet, providing an actionable candidate signal for fleet-wide correlation.
- **Endpoint Causal Context & Human Presence (T8, T19):** Flow-to-process binding, user session boundaries, interactive versus non-interactive process execution states, parent process lineage, service/task contexts, and correlated user hardware input telemetry (`GetLastInputInfo`, `WM_INPUT`, `/dev/input`). Auditing whether unprompted network bursts occur during deep workstation idle or screen-lock states.

- **Subsystem TLS Stack Consistency (T10, T20):** Verifying that processes generating port 443 HTTPS traffic legitimately load operating system TLS libraries (winhttp.dll, schannel.dll) and register corresponding ETW provider events, rather than establishing raw socket flows with disparate JA4 fingerprints.
- **Hardware Bus & Out-of-Band Boundaries (T16–T18):** Telemetry auditing across storage diagnostic pass-through routines (SCSI CDB 0x1D), IOMMU coherent DMA buffers, and hardware performance counters (PMU memory serialization barriers) to intercept non-network covert physical communication channels.
- **Process Execution Behavior:** Repetitive or quantized execution cycles observed in process runtime behavior, evaluated against workload-specific baselines as a candidate correlation signal.

No individual signal constitutes definitive proof; high-assurance detection decisions are reached through the correlation of multiple independent telemetry planes. The validation demanded by this approach must be conducted using workload-specific baselines and controlled adversary emulation simulations.

Operational Takeaway: When the discriminative power of flow-level network statistics diminishes, detection engineering shifts to the causality plane: the obsolete inquiry *"Does this network traffic appear human-like?"* is superseded by the causal question *"Is the host process, hardware bus state, and session context generating this communication channel consistent with the claimed user and workload identity?"*

Public-Release Scope Notice: This work does not disclose proprietary command-and-control implementation architectures or operational weaponization details. All examples and models are generalized exclusively to advance defensive engineering, telemetry measurement, and detection validation.

> **LEGAL AND INTELLECTUAL PROPERTY NOTICE (CC BY-NC-ND 4.0):**

> *"This research publication and technical specification is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0). It may not be used or reproduced without explicit author attribution; nor may it be integrated, directly or indirectly, into commercial cybersecurity products (EDR/NDR/XDR), managed services, or proprietary fee-based offerings without prior written authorization and an express commercial licensing agreement from the author. No derivative works or modifications are permitted. Document Identifier: SOC-NDR-MASTER-001-EN / Author: Oğulcan Erarslan (ByGhost)."*

PART II: COMPREHENSIVE TECHNICAL REPORT

> LEGAL AND INTELLECTUAL PROPERTY NOTICE (CC BY-NC-ND 4.0):

> "This research publication and technical specification is licensed under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International License (CC BY-NC-ND 4.0). It may not be used or reproduced without explicit author attribution; nor may it be integrated, directly or indirectly, into commercial cybersecurity products (EDR/NDR/XDR), managed services, or proprietary fee-based offerings without prior written authorization and an express commercial licensing agreement from the author. No derivative works or modifications are permitted. Document Identifier: SOC-NDR-MASTER-001-EN / Author: Oğulcan Erarslan (ByGhost)."

Executive Summary

This technical report evaluates measurement, modeling, and correlation methodologies available to enterprise defense against advanced Command-and-Control (C2) traffic. The operational scope is strictly confined to defensive detection: this document does not disclose adversary traffic generation algorithms, parameter estimation procedures, or evasion recipes (see Section 4 for scope boundaries). The technical evaluation addresses two structural questions:

1. Under what conditions do timing-based discriminators (jitter, periodicity, spectral line analysis) operate effectively, and under what conditions do they mathematically collapse?
2. When flow-level statistical separation fails, which telemetry planes continue to yield high-assurance defensive value?

Key Findings

- Jitter does not represent a single stochastic model; it encompasses at least two fundamentally divergent models that produce distinct Inter-Arrival Time (IAT) distributions and spectral signatures (Section 2.2). Conflating these models leads to erroneous theoretical justifications and invalid detection thresholds. This report deconstructs event-time jitter and interval jitter separately.
- **Under the uniform jitter assumption**, event-time jitter ($t_n = nT + \epsilon_n, \epsilon_n \sim U(-pT, pT)$) yields an IAT sequence $T + \epsilon_n - \epsilon_{(n-1)}$ following a symmetric triangular distribution with support $[T - 2pT, T + 2pT]$; conversely, interval jitter ($IAT_n = T + U_n, U_n \sim U(-pT, pT)$) produces an independent IAT sequence uniformly distributed over $[T - pT, T + pT]$. Dependency structures diverge fundamentally: event-time jitter exhibits an invariant lag-1 autocorrelation $\rho_1 = -0.5$ due to the shared ϵ noise term; interval jitter exhibits memoryless, independent intervals. Event-time generation mechanics differ: Model A defines $t_n = nT + \epsilon_n$ directly; Model B generates $t_n = \sum IAT_i$ cumulatively as a renewal process.
- Harmonic spectral lines originate strictly from the periodic point event train (the Dirac comb); the Discrete Fourier Transform (DFT) of a constant scalar IAT array yields exclusively a DC component ($f = 0$). Detection engines applying spectral analysis solely to scalar IAT sequences structurally miss fixed-period beacons.
- Exponential jitter shifts memorylessness not to the shifted variable $X = T + Y$, but strictly to the surplus variable Y ; the rigorous condition is $P(X > s + t \mid X > s) = e^{(-\lambda t)}$ for $s \geq T$ (Section 2.3).
- Extensive empirical traces across enterprise Ethernet, WAN, and cloud egress fabrics demonstrate self-similarity and Long-Range Dependence (LRD), exhibiting Hurst exponents in the persistent memory regime ($H > 0.5$). Poisson processes cannot reproduce this behavior; they flatten asymptotically to white noise under aggregation. This finding explains why variance and standard-deviation thresholds are fundamentally unstable across heavy-tailed workloads, necessitating tail-index and log-log diagnostic estimators (Section 3).
- The collapse of flow-level discrimination **is not an information-theoretic claim of absolute absence**. The claim applies strictly to the evaluated feature family: marginal IAT distributions, Power Spectral Density (PSD), autocorrelation decay, Hurst exponent, and diurnal profiles. Higher-order statistical dependencies, multifractal spectra, cross-flow multi-host correlation, and protocol state consistency remain viable discriminative surfaces (Section 5).

- The most resilient defensive measurement plane resides in the causal binding of network flow statistics with host process and session execution context. However, this plane is also conditional: session context divergence provides high-value signal in service and non-interactive execution models; processes executing inside interactive user sessions require deeper behavioral correlation (Sections 6.7–6.9).
- Twenty tactical implementations are cataloged (T1 – T20), expanding cross-layer causal binding from process/session context to human-presence verification, subsystem TLS consistency, and hardware/storage bus boundaries (Section 6).

Recommended Priority Actions

#	Action Item	Primary Owner	Time Horizon
1	Ingest moving-window scaling and tail-index estimators into flow analytics	Detection Engineering	0–30 Days
2	Bind process/session context to network flows via PID and ProcessStartTime	Endpoint Security	30–60 Days
3	Deploy timer quantization candidate feature testing calibrated by segment baseline	Network Security	30–60 Days
4	Conduct detector validation using an organization-controlled defensive profile simulator	Red / Blue Team	60–90 Days
5	Establish workload segmentation and quarterly baseline recalibration schedules	SOC Architecture	60–120 Days
6	Integrate maturity model tracking into quarterly executive security reviews	Security Leadership	90–365 Days

Executive Conclusion

At enterprise scale, C2 detection has fundamentally evolved from a perimeter network classification problem into a causal verification problem across intersecting network, endpoint, and identity telemetry planes. Capital investment must shift from tuning fragile flow-level thresholds to building cross-layer telemetry correlation pipelines.

1. Scope, Assumptions, and Threat Model

1.1 Scope

This technical report is engineered for Security Operations Centers (SOCs) defending mid-to-large enterprise infrastructures (5,000 to 100,000 endpoints). Scope topics include:

- The statistical foundations and failure boundaries of timing-based detection.
- Architectural blind spots in Network Detection and Response (NDR) flow-level feature spaces.
- Threat hunting tactics bridging network flow telemetry with host kernel instrumentation.
- Detection engineering deliverables: telemetry requirement matrices, Sigma rules, KQL queries, Splunk SPL, eBPF scripts, and maturity models.

Explicitly excluded: endpoint memory corruption exploitation, credential dumping mechanics, email initial access vectors, and adversary traffic generation code (see Section 4).

1.2 Assumptions and Constraints

- **Data Retention:** Flow telemetry (NetFlow v9, IPFIX, or connection logs) is retained for at least 90 days online and 12 months in cold archive. Required window lengths depend on sampling rates, workload stationarity, and seasonal variance (Section 8.4).
- **Temporal Resolution:** Connection-level summarizations (conn.log) are inadequate for sub-millisecond timing analysis; they record only start/end timestamps and byte aggregates. Sub-millisecond analysis requires full packet-level capture sensors on designated segments.
- **Transport Encryption:** The vast majority of enterprise egress traffic is TLS-encrypted. Detection does not rely on payload inspection, but on timing, volume, directionality, and process execution context.
- **Resource Budget (Design Target):** Analytic ingestion overhead is architecturally constrained to process less than 5% of daily flow volume in SIEM/data lake pipelines. This represents a capacity target, subject to local validation.
- **Standard of Proof:** This report strictly differentiates between candidate features (heuristics requiring correlation) and mathematical invariants.

1.3 Threat Model and MITRE ATT&CK Mapping

The threat model assumes sophisticated adversaries prioritizing low-and-slow persistence (long-dwell) over exfiltration velocity, shaping egress traffic to match organizational background baselines.

Observed Behavior	Compatible ATT&CK Technique	Architectural Note
Periodic egress via application protocols	T1071	Protocol choice alone lacks discriminative power
Encrypted communication channel	T1573.002	Ubiquitous across enterprise; requires contextual binding
Dynamic domain resolution infrastructure	T1568	Verified via passive DNS telemetry
Scheduled timing window transfers	T1029	Correlated against diurnal activity profiles
Protocol emulation / padding	T1001.003	Evaluated via JA4/JA3 and record size histograms
Timer-governed transmission cycles	T1497.003	Compatible with time-based evasion; subject to system confounders

2. Mathematical Foundations of Timing Analysis

2.1 Object of Detection: Event Train vs. IAT Sequence

A foundational error in network traffic analytics is conflating the point process of event timestamps (the event train) with the scalar sequence of inter-arrival times (IAT).

```

Event Train      :  $x(t) = \sum_n \delta(t - t_n)$ 
Fixed Period    :  $t_n = nT \Rightarrow X(f) = (1/T) \cdot \sum_k \delta(f - k/T)$ 
Scalar IAT Array :  $IAT_n = t_n - t_{(n-1)} = T \Rightarrow \text{DFT: Strictly DC component } (f = 0)$ 

```

The discrete harmonic lines spaced at $1/T$ in the spectrum originate from periodic event timestamps (via the Poisson summation formula). In contrast, the Discrete Fourier Transform (DFT) of a constant scalar IAT array yields solely a DC peak at bin zero ($X[0] = NT$), with all other bins identically zero.

- Assertions claiming "FFT of an IAT series exhibits harmonics at k/T " are mathematically false.
- The mathematically valid detection domain is the periodogram of the event train (or event-train autocorrelation), or the Lomb–Scargle periodogram.
- Operational impact: NDR products applying FFT directly to scalar IAT sequences structurally fail to detect constant-period beacons.

2.2 Two Distinct Jitter Models

Jitter does not represent a monolithic stochastic entity. The following two models generate distinct probability distributions, dependency structures, and spectral profiles:

Model A — Event-Time Jitter. Noise is added to absolute event timestamps:

$$t_n = nT + \varepsilon_n, \quad \varepsilon_n \text{ i.i.d.}, \quad E[\varepsilon] = 0, \quad \text{Var}(\varepsilon) = \sigma_\varepsilon^2$$

$$IAT_n = t_n - t_{(n-1)} = T + \varepsilon_n - \varepsilon_{(n-1)}$$

Mathematical properties:

- When $\varepsilon \sim U(-pT, pT)$, $IAT_n = T + \varepsilon_n - \varepsilon_{(n-1)}$ exhibits a **symmetric triangular distribution** with support $[T - 2pT, T + 2pT]$. Across general i.i.d. finite-variance noise, IAT_n remains symmetric around T , with $\text{Var}(IAT) = 2\sigma_\varepsilon^2$ and lag-1 autocorrelation $\rho_1 = -0.5$.
- **Chronological Ordering Condition:** For uniform bounded event-time jitter, $p < 0.5$ is a sufficient condition for guaranteed chronological ordering, because $T - 2pT > 0$ ensures every indexed inter-arrival interval is positive. When $p \geq 0.5$, chronological ordering is no longer guaranteed. Temporal reordering may occur, and indexed event order may cease to match physical time order.
- Adjacent intervals share the ε_n term: $\text{Cov}(IAT_n, IAT_{(n+1)}) = -\sigma_\varepsilon^2$, yielding **invariant lag-1 autocorrelation $\rho_1 = -0.5$** .
- **Spectral Structure (Displaced Lattice):** The Power Spectral Density (PSD) consists of discrete delta atoms at k/T scaled by $|\Phi_\varepsilon(k/T)|^2$ and a continuous background component scaled by $(1 - |\Phi_\varepsilon(f)|^2)$. For uniform jitter, $\Phi_\varepsilon(f) = \text{sinc}(2\pi pTf)$ (unnormalized sinc). Harmonic atoms are nullified when $2pk = m$ ($k = m/(2p)$, $m \in \mathbb{Z} \setminus \{0\}$), provided k coincides with an integer harmonic index.

Model B — Interval Jitter. Noise is added directly to interval durations:

$$IAT_n = T + U_n, \quad U_n \text{ i.i.d.}, \quad E[U] = 0, \quad \text{Var}(U) = \sigma_U^2$$

Uniform Example : $U_n \sim \text{Uniform}(-pT, pT)$

$$t_n = \sum_{i \leq n} IAT_i \quad (\text{Cumulative Renewal Process})$$

Mathematical properties:

- In the uniform case, IAT_n is independent and uniformly distributed over $[T - pT, T + pT]$.
- **Positivity Condition:** Valid renewal processes require strictly positive intervals: $P(T + U_n > 0) = 1$, requiring **$p < 1$** .
- Event times accumulate independent increments: phase error executes a random walk.
- **Spectral Structure:** For non-lattice continuous interval distributions, the spectrum for $f \neq 0$ is **strictly continuous; discrete line components do not exist**. Continuous PSD is given by the renewal formula: $S(f) = \lambda \cdot (1 - |\psi(f)|^2) / |1 - \psi(f)|^2$, where $\lambda = 1/T$ and $\psi(f)$ is the characteristic function of the interval distribution.

Critical Architectural Distinction: Model A contains discrete harmonic atoms alongside a continuous component; Model B contains strictly continuous power for $f \neq 0$. Coupling the IAT distribution of Model A with the spectral profile of Model B is mathematically invalid.

2.3 Exponential Jitter and the Scope of Memorylessness

Exponential jitter is not a zero-mean perturbation: for $Y \sim \text{Exp}(\lambda)$, $E[Y] = 1/\lambda$ and $\text{Var}(Y) = 1/\lambda^2$. Correct formulation:

$$X = T + Y, \quad Y \sim \text{Exp}(\lambda), \quad \text{Support}(X) = [T, \infty)$$

Surplus Variable : $Y = X - T$

Memorylessness : $P(Y > y + t \mid Y > y) = P(Y > t) = e^{(-\lambda t)}$

Equivalent Form : $P(X > s + t \mid X > s) = e^{(-\lambda t)}, \quad s \geq T$

Memorylessness belongs strictly to the surplus variable Y , not to shifted variable X . The expression $P(X > s + t \mid X > s) = P(X > t)$ is mathematically false. Furthermore, the sharp lower cutoff at T remains an observable candidate anomaly.

2.4 Sinc Conventions

This technical report adheres strictly to the unnormalized sinc convention:

Unnormalized (Used Here) : $\text{sinc}(x) \equiv \sin(x) / x$, $\text{sinc}(0) = 1$
 Normalized (Not Used) : $\text{sinc}_N(x) \equiv \sin(\pi x) / (\pi x)$

The spectral envelope is expressed as $\Phi(f) = \text{sinc}(2\pi f T_f)$.

2.5 Bounding the Detection Ceiling Claim

This report does not assert that shaped C2 traffic becomes undetectable in an absolute information-theoretic sense. The detection ceiling claim is strictly bounded:

> Within the evaluated flow-level feature family—marginal IAT distributions, Power Spectral Density, autocorrelation, Hurst exponent, and diurnal cycles—cross-class discrimination between shaped C2 and legitimate background traffic degrades drastically.

Discriminative capacity is preserved across orthogonal feature spaces:

- Higher-order statistical dependencies (bispectra, polyspectral entropy),
- Multifractal spectrum widths ($h(q)$),
- Cross-flow multi-host correlation,
- Application/protocol state consistency (TLS record sizes, JA4 fingerprints),
- Causal host process and session execution telemetry.

2.6 Superposition and Asymptotic Regimes (Palm-Khintchine)

Convergence of aggregated point processes to a Poisson process under the Palm-Khintchine theorem requires three strict conditions:

1. **Mutual Independence** between constituent sources.
2. **Absence of Dominant Sources:** No single source contributes a disproportionate fraction of aggregate event rate.
3. **Rare-Source Regime:** Individual source rates remain infinitesimal relative to aggregate throughput.

When these conditions fail (e.g., enterprise WAN links dominated by heavy-tailed transfers), aggregate traffic preserves non-Poisson burstiness and self-similarity.

2.7 Limitations of Machine Learning-Based Detection

Unsupervised autoencoders learn normal representations and trigger anomalies based on reconstruction error. However, their discriminative performance is completely governed by input feature representations. When Model A timing is applied, the triangular IAT distribution and $\rho_1 = -0.5$ autocorrelation provide candidate features; under Model B, discrete line features vanish, forcing anomaly models to rely on continuous spectral shapes and tail cutoffs. Claiming that deep neural networks inherently solve timing evasion without rigorous feature engineering is invalid.

3. Statistical Fingerprint of Enterprise Traffic: Self-Similarity & Long-Range Dependence

3.1 Scaling Laws and the Hurst Exponent

Seminal studies (Leland et al. 1994, Paxson & Floyd 1995) established that enterprise Ethernet and WAN traffic exhibit self-similarity and Long-Range Dependence (LRD), characterized by Hurst parameters $0.5 < H < 1.0$.

Under temporal aggregation by factor m , a discrete process X behaves as:

$$X^{(m)}[k] = (1/m) \cdot \sum_{i=(k-1)m+1}^{km} X[i]$$

$$\text{Var}(X^{(m)}) \sim \sigma^2 \cdot m^{(2H-2)}$$

While Poisson processes decay rapidly ($\text{Var} \sim 1/m$), LRD traffic exhibits hyperbolic memory retention. Detectors evaluate aggregated count series (stationary increments resembling fractional Gaussian noise, fGn), distinct from non-stationary cumulative processes (fractional Brownian motion, fBm).

3.2 Heavy Tails and the Source Model

Willinger et al. demonstrated that aggregating independent ON/OFF sources with heavy-tailed durations ($P(\text{ON} > x) \sim x^{-(\alpha)}, 1 < \alpha < 2$) provably generates LRD traffic in the limit:

$$H = (3 - \alpha) / 2$$

In the regime $1 < \alpha < 2$, source phase durations exhibit infinite variance. Consequently, sample variances computed over phase durations fail to converge, rendering classical variance and standard deviation thresholds unstable.

K-point Hill estimators evaluate regular tail variation:

$$\hat{\alpha}_{\text{Hill}}(k) = \left[\frac{1}{k} \cdot \sum_{i=1}^k \ln(x_{(i)} / x_{(k+1)}) \right]^{-1}$$

Log-log Complementary Cumulative Distribution Function (CCDF) plots serve as diagnostic tools to identify linear scaling regions before calibrating Hill thresholds.

3.3 Multifractality

Enterprise traffic traces often exhibit multifractal scaling, where moments of order q scale non-linearly with exponent $h(q)$. Monofractal synthetic traffic generated by single-parameter algorithms exhibits an abnormally narrow $h(q)$ spectrum width, providing an orthogonal candidate feature.

3.4 Defensive Implications

Self-similarity provides two structural opportunities to Blue Teams:

1. Scaling parameter estimations stabilize over long observation windows.
2. Scaling parameters reflect aggregate host behavior across all egress flows, allowing host-level profiling.

4. Scope Boundary and Non-Operational Model Context

4.1 Disclosure Policy

This report focuses strictly on defensive detection. It intentionally excludes deployable generator code, target-environment parameter extraction scripts, and operational evasion playbooks.

4.2 Model Taxonomy

Model Class	Sojourn / Phase Duration	Autocorrelation Decay	Typical Scaling	Impact on Detector Assumptions
Classical MMPP	Exponential	Short-Range (SRD)	$H \approx 0.5$	Exponential decay; rapid convergence of window means
Semi-Markov Modulated	Heavy-Tailed	Long-Range (LRD)	$H > 0.5$	Infinite phase variance; requires tail estimators
Pareto ON/OFF Superposition	Heavy-Tailed	LRD (fBm limit)	$H = (3 - \alpha)/2$	Scaling highly sensitive to observation window

4.3 Distinction Between fBm and fGn

- **fBm (Fractional Brownian Motion):** Continuous, non-stationary cumulative process with stationary increments. Does not possess a classical PSD.
- **fGn (Fractional Gaussian Noise):** Stationary increment process. Low-frequency power spectrum behaves as $S(f) \sim |f|^{-(1-2H)}$.

4.4 Mapping Detector Assumptions to Regimes

Detector Family	Underlying Model Assumption	Invalidation Regime
Fixed Period / FFT Peak	Discrete spectral lines in point process	Transition to continuous renewal PSD
IAT Variance Threshold	Finite, converging second moment	Heavy-tailed sojourn times ($1 < \alpha < 2$)
Autocorrelation Periodicity	Short-range, periodic dependence	LRD regime (hyperbolic memory decay)
Window Mean Comparison	Approximate sample independence	LRD regime (degrees-of-freedom distortion)
Marginal Distribution Matching	Cross-class distributional divergence	Baseline kernel density matching
Entropy Anomaly Detector	Divergent distributional uncertainty	Equifinal distributions sharing identical entropy

5. The Scope of the Detection Ceiling

5.1 Status of the Flow-Level Feature Space

Feature Group	Fixed-Rate Beacon	Jittered Beacon (Model A / B)	Baseline-Adapted Shaped C2
Spectral Peak	Sharp, prominent	Attenuated lines (A) / strictly continuous (B)	Absent from evaluated family
Marginal IAT Distribution	Deterministic spike	Triangular (A) / uniform (B)	Matches baseline kernel density
Autocorrelation	Strictly periodic	$\rho_1 = -0.5$ (A) / independent (B)	Hyperbolic decay (LRD)
Scaling Exponent (H)	Undefined	Short-range ($H \approx 0.5$)	Matches enterprise subnet band
PSD Envelope	Discrete delta comb	Modulated sinc atoms (A) / smooth renewal PSD (B)	Continuous $1/f^\beta$ curve
Diurnal Cycle	None	None	Synchronized with business hours

5.2 Structural Fallacies of Classical Detectors

Single-dimensional threshold rules ($\text{std}(\text{IAT}) > \sigma$) generate intolerable false positive volumes on bursty enterprise links. Periodicity detectors fail against continuous renewal spectra. Entropy metrics fail because fundamentally divergent distributions produce identical scalar Shannon entropy values.

5.3 The Labeling Dilemma and Concept Drift

Supervised classifiers trained on public C2 datasets overfit to obsolete, rigid beacon profiles. Enterprise baseline drift (cloud migrations, remote work shifts) rapidly invalidates static machine learning models.

5.4 Remaining Discriminative Surfaces

Defense-in-depth leverages:

- High-order statistical dependencies (bispectra),
- Multifractal spectrum widths ($h(q)$),
- Cross-flow multi-host synchronization,
- TLS fingerprinting (JA4, record size distributions),

- Timer quantization artifacts,
- Endpoint causal context (flow-to-process binding).

6. Tactical Catalog (T1 – T20)

Every tactic is cataloged across nine formal engineering specifications:

6.1 T1 — Moving-Window Scaling Stability

- **Surface:** Network.
- **Data Source:** NetFlow v9 / IPFIX / connection logs; byte counts or connection counts per window.
- **Hypothesis:** On human-driven endpoints, scaling exponents fluctuate across days; algorithmic traffic maintains abnormal parameter stability.
- **Rationale:** Under LRD, external perturbations induce parameter variance. Static algorithmic generators exhibit unnaturally low variance across sliding estimation windows.
- **Confounders:** Machine-to-machine sync agents, telemetry beacons, weekend work shifts.
- **False Positives:** Automated cloud backup daemons, RMM management agents.
- **Required Baseline:** Workload-specific segment baseline collected over at least 90 days.
- **Validation:** Controlled injection of synthetic constant-profile traffic (T15).
- **Evidentiary Weight:** Requires Correlation.

6.2 T2 — Tail Structure and Hill Estimator Deviation

- **Surface:** Network.
- **Data Source:** Flow/session byte volumes and burst durations.
- **Hypothesis:** Authentic heavy-tailed traffic exhibits smooth Hill plots; artificial truncation introduces abrupt parameter deviation beyond threshold k .
- **Rationale:** Hard limits on C2 sleep intervals or burst buffers truncate empirical CCDF tails.
- **Confounders:** TLS record limits, MTU boundaries, application buffer caps.
- **False Positives:** Video streaming, bulk file sync.
- **Required Baseline:** Application/protocol-specific CCDF baselines.
- **Validation:** Benchmarking against synthetic truncated distributions.
- **Evidentiary Weight:** Requires Correlation.

6.3 T3 — Quantization and Comb Pattern (Candidate Feature)

- **Surface:** Network (Packet Level).
- **Data Source:** High-precision packet timestamps (TAP/SPAN capture).
- **Hypothesis:** Software timers introduce clustering at discrete OS scheduling quantum multiples.
- **Rationale:** Non-real-time OS schedulers (e.g., Windows 15.6ms quantum) quantize thread wakeups.
- **Confounders:** TCP pacing, NIC interrupt coalescing, VM hypervisor clock drift, NTP adjustments.
- **False Positives:** Media streaming, periodic polling agents.
- **Required Baseline:** Sensor clock accuracy certification and segment phase histograms.
- **Validation:** Testing against known timer-driven internal services.
- **Evidentiary Weight:** Requires Correlation.

6.4 T4 — Fleet-Wide Common Parameter Clustering

- **Surface:** Network + Identity.
- **Data Source:** Host-level scaling and tail summaries enriched by directory inventory.
- **Hypothesis:** Disparate hosts in different roles exhibiting statistically identical timing parameters indicate common automated orchestration.
- **Rationale:** Human workloads diverge across roles and geographies.
- **Confounders:** Standard OS base images, corporate software update agents.
- **False Positives:** Identical EDR/RMM agents executing on uniform hardware.

- **Required Baseline:** Fleet-wide parameter distribution across role clusters.
- **Validation:** Benchmarking against known fleet-wide software rollouts.
- **Evidentiary Weight:** Requires Correlation.

6.5 T5 — Regime Shifts and Structural Break Analysis

- **Surface:** Network.
- **Data Source:** Host-level time series; windowed PSD and autocorrelation summaries.
- **Hypothesis:** Legitimate traffic experiences frequent external breaks; absence of expected structural breaks serves as a candidate anomaly.
- **Rationale:** Classical break tests (CUSUM, Bai-Perron) must be calibrated for LRD memory retention to avoid inflated rejection rates.
- **Confounders:** Seasonality, maintenance windows, holiday shifts.
- **False Positives:** Continuous background data ingest services.
- **Required Baseline:** At least two quarters of seasonal history.
- **Validation:** Correlation against change management tickets.
- **Evidentiary Weight:** Standalone (with contextual enrichment).

6.6 T6 — Endpoint Causal Correlation (ETW / eBPF)

- **Surface:** Endpoint + Network.
- **Data Source:** Windows ETW network/process providers; Linux eBPF tracepoints (`sock:inet_sock_set_state`).
- **Hypothesis:** Correlating flow statistical profiles with initiating process lineage uncovers anomalous execution contexts.
- **Rationale:** Statistical flow properties and host execution contexts represent independent observation planes.
- **Confounders:** PID reuse, container network namespaces, telemetry collection latency.
- **False Positives:** Browser background rendering tabs, update helpers.
- **Required Baseline:** Process inventory and normal network usage mapping.
- **Validation:** End-to-end verification via synthetic PID reuse tests.
- **Evidentiary Weight:** Correlation Infrastructure.

6.7 T7 — Session and Privilege Context Validation

- **Surface:** Endpoint + Identity.
- **Data Source:** Logon session ID, desktop attachment, token privileges, parent-child process tree.
- **Hypothesis:** Headless, non-interactive processes generating continuous external communications represent high-value candidate anomalies.
- **Rationale:** Non-interactive execution boundaries provide high signal-to-noise differentiation.
- **Confounders:** Legitimate service accounts, scheduled maintenance tasks.
- **False Positives:** CI/CD runners, enterprise management services.
- **Required Baseline:** Service account network communication profiles.
- **Validation:** Inventory-based exception modeling.
- **Evidentiary Weight:** Requires Correlation.

6.8 T8 — Input Event and Network Activity Correlation

- **Surface:** Endpoint.
- **Data Source:** HID input event telemetry (`WM_INPUT`, `/dev/input`), window focus state.
- **Hypothesis:** Interactive workstation endpoints generating bursty egress in the complete absence of user input events represent candidate automated processes.
- **Rationale:** Human-driven applications correlate probabilistically with physical keystroke/mouse activity.
- **Confounders:** Kiosk workstations, background audio streaming, tab sync.
- **False Positives:** Cloud file sync, automated auto-save routines.
- **Required Baseline:** Device-class input/network correlation distributions.
- **Validation:** Profiling verified background sync workloads.

- **Evidentiary Weight:** Requires Correlation.

6.9 T9 — Timer and Sleep Pattern Telemetry

- **Surface:** Endpoint.
- **Data Source:** Thread sleep API hooks (ntd11!NtDelayExecution, kernel32!SleepEx), CPU time distribution.
- **Hypothesis:** Process thread sleep distributions diverge from expected workload baselines.
- **Rationale:** Programmed sleep patterns leave measurable thread state artifacts.
- **Confounders:** Queue consumers, exponential backoff retries.
- **False Positives:** Database connection pool managers, health monitoring daemons.
- **Required Baseline:** Application-specific thread sleep baselines.
- **Validation:** Benchmarking verified enterprise agent thread traces.
- **Evidentiary Weight:** Requires Correlation.

6.10 T10 — TLS and Application Layer Consistency

- **Surface:** Network.
- **Data Source:** Passive TLS fingerprinting (JA4, cipher suites, ALPN), record size distributions.
- **Hypothesis:** Discrepancies exist between declared user-agent identities and empirical TLS stack negotiation signatures.
- **Rationale:** Emulating complete client protocol stacks requires substantial engineering beyond timing shaping.
- **Confounders:** Corporate SSL inspection proxies, browser auto-updates.
- **False Positives:** Outdated embedded clients, custom API integrations.
- **Required Baseline:** Application-specific JA4 fingerprint catalogs.
- **Validation:** Correlation against software deployment inventories.
- **Evidentiary Weight:** Requires Correlation.

6.11 T11 — Packet Size Distribution and Volume Clustering

- **Surface:** Network.
- **Data Source:** Packet length histograms per flow.
- **Hypothesis:** Multimodal legitimate application distributions contrasting with unimodal or unnaturally clustered distributions indicate automated telemetry.
- **Rationale:** Real application protocols generate diverse frame sizes based on data structures.
- **Confounders:** Transport compression, encryption padding, MTU limits.
- **False Positives:** Telemetry forwarders, IoT sensors.
- **Required Baseline:** Protocol-class packet size histograms.
- **Validation:** Comparison against known protocol reference traces.
- **Evidentiary Weight:** Requires Correlation.

6.12 T12 — Cross-Layer Correlation Architecture

- **Surface:** Network + Endpoint + Identity.
- **Data Source:** Composite join key: (host_id, pid, process_start_time, 5_tuple, event_time).
- **Hypothesis:** Combining ambiguous network observations with endpoint and identity contexts produces high-confidence verdicts.
- **Rationale:** Multi-planar telemetry correlation eliminates single-sensor ambiguity. Requiring ProcessStartTime prevents PID-reuse collision.
- **Confounders:** Time synchronization drift, telemetry pipeline drops.
- **False Positives:** Dropped event joins resulting in phantom alerts.
- **Required Baseline:** Telemetry pipeline health and join success metrics.
- **Validation:** Synthetic test flows validating end-to-end join integrity.
- **Evidentiary Weight:** Correlation Infrastructure.

6.13 T13 — Deception and Canary Flows

- **Surface:** Network.
- **Data Source:** Monitored decoy services, canary DNS tokens, and honeypot egress endpoints.
- **Hypothesis:** Interaction with unadvertised canary assets cannot be justified by authentic business activity.
- **Rationale:** Adversary recon and profile-matching routines interact with decoy lures.
- **Confounders:** Misconfigured internal vulnerability scanners.
- **False Positives:** Internal security audits, crawler scripts.
- **Required Baseline:** Decoy asset registry and exclusion policies.
- **Validation:** Periodic automated canary verification probes.
- **Evidentiary Weight:** Standalone.

6.14 T14 — Long-Horizon Monitoring and Seasonality

- **Surface:** Network.
- **Data Source:** 90–365 days of archived flow aggregates.
- **Hypothesis:** Authentic enterprise workloads exhibit seasonal cycles (quarter-end surges, holiday dips); artificial traffic fails to adapt.
- **Rationale:** Long observation horizons suppress estimator variance and isolate seasonal components.
- **Confounders:** Organizational restructuring, sudden business shifts.
- **False Positives:** Ad-hoc project workloads.
- **Required Baseline:** At least one full annual operational cycle.
- **Validation:** Hold-out forecasting validation.
- **Evidentiary Weight:** Requires Correlation.

6.15 T15 — Detector Validation via Defensive Profile Simulator

- **Surface:** All Planes.
- **Data Source:** Isolated test harness executing controlled, benign-like synthetic traffic.
- **Hypothesis:** Detector response curves and false positive thresholds can be empirically mapped using controlled traffic generators.
- **Rationale:** Uncalibrated detectors exhibit unquantified operational efficacy.
- **Confounders:** Testbed network divergence from production fabrics.
- **False Positives:** Simulation traffic leaking into production monitoring.
- **Required Baseline:** Production segment baseline vs. simulator profile.
- **Validation:** Sensitivity curve mapping and marginal contribution scoring (Section 8.3).
- **Evidentiary Weight:** Validation Infrastructure.

6.16 T16 — Storage Subsystem Pass-Through Auditing

- **Surface:** Endpoint (Storage / Block I/O Layer).
- **Data Source:** Windows: Microsoft-Windows-StorPort ETW provider (EventID 500-502), DeviceEvents (DeviceIoControl). Linux: eBPF tracepoint/syscalls/sys_enter_ioctl (SG_IO 0x2285), tracepoint/block/block_rq_issue.
- **Hypothesis:** Non-storage user-mode processes issuing direct diagnostic SCSI command blocks (CDB Opcode 0x1D - SEND DIAGNOSTIC or vendor-specific opcodes 0xC0-0xFF) to raw physical disks (\\.\PhysicalDrive*, /dev/sd*) indicate covert out-of-band communication channels.
- **Rationale:** Traditional network telemetry stacks (WFP, NDIS, eBPF TC) do not monitor block storage subsystems. Tunneling data through storage diagnostic passes bypasses network-layer detection.
- **Confounders:** RAID/NVMe management suites (StorCLI, MegaRAID, Intel RST), disk SMART monitoring daemons (smartctl).
- **False Positives:** Hardware vendor diagnostics, raw volume snapshot backup agents.
- **Required Baseline:** Enterprise inventory of binaries authorized for raw block device I/O.
- **Validation:** Controlled injection of IOCTL_SCSI_PASS_THROUGH_DIRECT diagnostic frames.

- **Evidentiary Weight:** Standalone (High-Priority Triage).

6.17 T17 — IOMMU / DMA Coherent Buffer & Ring Descriptor Auditing

- **Surface:** Hardware / Kernel (Bus Layer).
- **Data Source:** Kernel DMA Protection (DMAR) audit logs, hypervisor IOMMU page faults, kernel allocation telemetry (`dma_alloc_coherent`, `AllocateCommonBuffer`).
- **Hypothesis:** Unregistered or non-network PCIe endpoints allocating high-frequency circular ring buffers in host DRAM and triggering MMIO doorbells represent candidate hardware egress bridges (e.g., SmartNIC, FPGA).
- **Rationale:** Enterprise DMA channels are mediated by IOMMU. Physical buffer allocation patterns diverge between standard compute/GPU workloads and continuous communication rings.
- **Confounders:** High-performance GPU computing workloads (CUDA/OpenCL), high-throughput NVMe queue management.
- **False Positives:** Machine learning accelerator cards, high-speed frame grabbers.
- **Required Baseline:** Enterprise PCIe hardware inventory (Vendor/Device IDs) and IOMMU mapping profiles.
- **Validation:** Controlled emulation of ring buffer allocations via test PCIe endpoints.
- **Evidentiary Weight:** Requires Correlation.

6.18 T18 — Hardware Performance Counter (PMU) & Memory Barrier Frequency

- **Surface:** Hardware / CPU Microarchitecture.
- **Data Source:** Hardware Performance Monitoring Units (PMU via `perf_event_open` on Linux, ETW PMU events on Windows).
- **Hypothesis:** Processes executing abnormal frequencies of memory serialization barriers (SFENCE, LFENCE, MFENCE) paired with direct writes to MMIO memory-mapped regions deviate from standard application execution patterns.
- **Rationale:** Writing coherent frames to hardware DMA rings requires explicit CPU store-fence synchronization to enforce memory ordering before doorbell ringing.
- **Confounders:** High-Frequency Trading (HFT) algorithms, lock-free multithreading libraries, graphics rendering engines.
- **False Positives:** Concurrent runtime engines utilizing atomic operations.
- **Required Baseline:** Process-class PMU instruction mix and memory fence baselines.
- **Validation:** Microarchitectural profiling against reference lock-free algorithms.
- **Evidentiary Weight:** Candidate Feature (Requires Correlation).

6.19 T19 — Human-Presence & HID Interaction Causality Correlation

- **Surface:** Endpoint + Network (Cross-Layer Causality).
- **Data Source:** Windows: `GetLastInputInfo`, `WM_INPUT`, workstation lock state events. Linux: `/dev/input/*`, `evdev` events, session idle counters. Network: Egress NetFlow / socket byte counters.
- **Hypothesis:** Workstation processes executing in interactive user sessions (Session 1+) that generate periodic or bursty outbound network traffic during prolonged user inactivity (idle > 60s) or locked desktop states represent candidate automated agents.
- **Rationale:** Authentic interactive applications exhibit causal dependence on human input activity ($P(\text{Egress}_t > 0 \mid \text{recent HID activity}) \neq P(\text{Egress}_t > 0 \mid \text{idle/locked state})$). Automated channels lack temporal coupling with physical keystroke and mouse inputs.
- **Confounders:** Automatic browser tab refreshers, active audio/video streaming in background tabs.
- **False Positives:** Enterprise background backup or signature update jobs initiated upon user lock.
- **Required Baseline:** Device-type classification (Workstation vs. Server) and application-class HID correlation matrices.
- **Validation:** Profiling network egress behavior during verified idle and lock intervals.
- **Evidentiary Weight:** High Correlation Power.

6.20 T20 — Subsystem TLS Stack Divergence & ETW WinHTTP/Schannel Verification

- **Surface:** Endpoint (Module Execution) + Network (TLS Handshake).
- **Data Source:** Windows: Microsoft-Windows-Kernel-Process (ImageLoad), Microsoft-Windows-WinHttp ETW provider, Virtual Address Descriptors (VAD). Network: Port 443 TCP flows and passive JA4 fingerprints.
- **Hypothesis:** Processes establishing outbound TCP port 443 connections that lack operating system HTTP/TLS modules (winhttp.dll, schannel.dll) in their loaded module lists and emit zero Microsoft-Windows-WinHttp ETW events represent standalone or custom cryptographic engines.
- **Rationale:** Authentic enterprise desktop applications leverage operating system HTTP/TLS stacks to inherit corporate proxy settings, certificate validation, and centralized audit logging. Disparate TLS implementations produce JA4 signature divergence and module absence.
- **Confounders:** Legitimate developer runtimes (Go, Rust, Python, Node.js, curl) utilizing statically compiled TLS libraries.
- **False Positives:** Electron applications utilizing embedded BoringSSL engines.
- **Required Baseline:** Inventory of sanctioned enterprise binaries utilizing custom or embedded TLS runtimes.
- **Validation:** Comparing outbound socket creation with ETW module load and WinHTTP trace telemetry.
- **Evidentiary Weight:** Requires Correlation (Complements T6 and T10).

7. Detection Engineering Deliverables

7.1 Telemetry Requirement Matrix and Provider Notes

Telemetry Stream	Source / Provider	Target Tactics	Nominal Granularity	Priority
Flow Records	IPFIX / NetFlow v9 / Zeek	T1, T2, T4, T5, T14	1 Second	Critical
Packet Timestamps	TAP / SPAN hardware sensors	T3, T11	Microsecond	High
Process Execution Events	Windows ETW (Microsoft-Windows-Kernel-Process)	T6, T7	100 Nanoseconds	Critical
Flow-to-PID Mapping	Windows ETW (Microsoft-Windows-Kernel-Network)	T6, T12	Event-Driven	Critical
Linux Kernel Sockets	eBPF tracepoints (inet_sock_set_state)	T6, T7, T9	Nanosecond	Critical
User Input Telemetry	Endpoint Agent (WM_INPUT, /dev/input)	T8, T19	Millisecond	High
TLS Fingerprints	Passive TLS Sensors (JA4 / JA3)	T10, T20	Session-Driven	High
Storage Pass-Through I/O	Windows StorPort ETW / Linux eBPF SG_IO	T16	Event-Driven	High
Kernel DMA / IOMMU Logs	DMAR Audit Logs / IOMMU Page Faults	T17	Event-Driven	High
CPU Performance Counters	Hardware PMU / perf_event_open	T18	Event-Driven	Medium
WinHTTP / Schannel Calls	Windows ETW (Microsoft-Windows-WinHttp)	T20	Event-Driven	High
Directory & Identity	SIEM / IdP / Kerberos logs	T7, T12	Event-Driven	High

Resolution Caution: Nominal timestamp units (e.g., 100 ns) do not represent effective measurement resolution. True resolution is dictated by clock sources (TSC vs. HPET), interrupt latencies, and NTP/PTP discipline.

7.2 Process Context Rule (Sigma, Reference)

```

title: Hidden or Non-Interactive Script Host Launched by Scheduler
id: 8f3c1a72-4d5e-4a91-9a1e-7c2b6d0e5f31
status: experimental
logsource:
  category: process_creation
  product: windows
detection:
  selection_parent:
    ParentImage|endswith:
      - '\svchost.exe'
      - '\taskeng.exe'
      - '\schtasks.exe'
  selection_child:
    Image|endswith:
      - '\powershell.exe'
      - '\pwsh.exe'
      - '\wscript.exe'
      - '\python.exe'
  selection_flags:
    CommandLine|contains:
      - '-WindowStyle Hidden'
      - '-NonInteractive'
      - '-EncodedCommand'
  condition: selection_parent and selection_child and selection_flags
fields:
  - Computer
  - User
  - ParentImage
  - Image
  - CommandLine
falsepositives:
  - Enterprise RMM and management agents
  - System configuration management scripts (SCCM / Intune)
level: medium

```

7.3 Scaling and Context Correlation (KQL, Reference)

```

// Reference KQL – Schemas and field names must be adapted to local data lake definitions.
let lookback = 30d;
let scaling = TrafficScalingMetrics_CL
| where TimeGenerated > ago(lookback) and SampleCount > 50
| summarize H_avg = avg(HurstWavelet), H_std = stdev(HurstWavelet), Samples = sum(SampleCount) by DeviceId
// H_avg range is an illustrative baseline; derive thresholds from segment baselines (Section 8.1).
| where Samples > 2000 and H_avg between (0.55 .. 0.90) and H_std < 0.05;
let context = EndpointContext_CL
| where TimeGenerated > ago(lookback) and IsInteractive == false
| summarize HeadlessCount = dcount(ProcessId), Examples = make_set(SessionId, 5)
  by DeviceId, ProcessId, ProcessStartTime;
scaling
| join kind=inner context on DeviceId
| project DeviceId, H_avg, H_std, Samples, HeadlessCount, ProcessStartTime
| order by H_std asc

```

7.4 Quantization Candidate Feature Test (Splunk SPL & Python)

```

index=sensor sourcetype="c2sensor:packet"
| sort 0 + _time
| streamstats current=f last(_time) as prev_ts by flow_id
| eval iat_ms = round((_time - prev_ts) * 1000, 3)
| where isnotnull(iat_ms) AND iat_ms > 0
| eval phase = iat_ms % q
| eval edge = min(phase, q - phase)
| stats count as total, sum(eval(if(edge < tau, 1, 0))) as edge_hits by src_ip, dest_ip
| eval edge_share = round(edge_hits / total, 4)
| where total > 500 AND edge_share > theta_edge
| table src_ip, dest_ip, total, edge_share

```

```
import numpy as np
```

```
def quantum_comb_score(ts_ns: np.ndarray, q_ms: float, tau_ms: float,
                      theta_edge: float, theta_chi2: float) -> dict:
```

```

\\\\"Evaluates timer quantization clustering on timestamp series (ns).\\\\"
ts = np.sort(np.asarray(ts_ns, dtype=np.int64))
iat_ms = np.diff(ts) / 1e6
iat_ms = iat_ms[iat_ms > 0]
if iat_ms.size < 500:
    return {"n": int(iat_ms.size), "verdict": "insufficient_sample"}

phase = np.mod(iat_ms, q_ms)
edge = np.minimum(phase, q_ms - phase)
edge_share = float(np.mean(edge < tau_ms))

counts, _ = np.histogram(phase, bins=16, range=(0.0, q_ms))
expected = counts.sum() / 16.0
chi2 = float(np.sum((counts - expected) ** 2 / expected))

return {
    "n": int(iat_ms.size),
    "edge_share": round(edge_share, 4),
    "chi2": round(chi2, 2),
    "verdict": "candidate" if edge_share > theta_edge and chi2 > theta_chi2 else "inconclusive",
}

```

7.5 Process Context Instrumentation (eBPF bpftrace & libbpf CO-RE)

```

/* bpftrace: Byte volume per process in 60s windows
tcp_sendmsg(struct sock *sk, struct msghdr *msg, size_t size) -> arg2 = size */
kprobe:tcp_sendmsg
{
    @bytes[comm, pid] = sum(arg2);
}
interval:s:60
{
    print(@bytes);
    clear(@bytes);
}

```

```

/* libbpf CO-RE: Identify processes lacking a controlling terminal */
static __always_inline int has_controlling_tty(void)
{
    struct task_struct *task = (struct task_struct *)bpf_get_current_task_btf();
    struct signal_struct *sig = BPF_CORE_READ(task, signal);

    if (!sig)
        return 0;
    return BPF_CORE_READ(sig, tty) != NULL;
}

```

7.6 Detector Lifecycle Management

Every production detector must maintain explicit metadata: unique ID, owner, telemetry dependencies, rule logic version, documented false positive sources, evidentiary class, and last validation date. Detectors uncalibrated for over 90 days are marked "unverified."

8. False Positive Management and Quantitative Metrics

8.1 Workload Segmentation and Qualitative Baseline

Segment	Scaling Behavior	Diurnal Profile	Architectural Note
Knowledge Worker Desktops	Strong LRD ($H > 0.7$)	Sharp business-hour curve	Highest priority hunt segment
Developer Workstations	Moderate–Strong LRD	High off-hours activity	High false positive risk
Application Servers	Strong LRD	Weak diurnal dependency	High traffic continuity
Backup / Bulk Ingestion	Dominated by bursts	Window-constrained	Requires dedicated exclusion
Kiosk / Embedded Ends	Weak–Moderate LRD	Strict business window	Limited user input telemetry

8.2 Sources of False Positives

- **Enterprise Cloud Storage:** Large file synchronizations mimic heavy-tailed bursts.
- **Content Delivery Networks (CDNs):** Timing governed by edge infrastructure.
- **RMM and EDR Management Agents:** Highly regular, low-volume heartbeats.
- **CI/CD Automation Pipelines:** Non-human, bursty egress.

8.3 Metrics Framework

- **Signal-to-Noise Ratio (SNR):** Verified True Positives per total alert volume.
- **Marginal Contribution:** Incremental detection rate established via simulation (T15).
- **Telemetry Health Score:** Data drop rates, collection latencies, and schema completion rates.
- **Mean Time to Detect (MTTD):** Measured across quarterly adversary emulation drills.

8.4 Determination of Window Lengths

Observation windows (30, 90, 180 days) are operational baselines. Sizing depends on sampling rates, process stationarity, estimator bias-variance trade-offs, and computational storage budgets.

8.5 The Headless Server Paradox and Dynamic Baseline Partitioning

A critical architectural pitfall in cross-layer correlation is the naive assumption of uniform endpoint telemetry across workstations and servers. Correlation rules binding network egress to interactive human input (T8, T19) generate massive false-positive storms when deployed across application servers, container nodes, or virtualization hosts where zero HID interaction naturally occurs.

Defensive frameworks must partition causality rules dynamically:

- **Interactive Workstation Profile:** Enforces strict human-presence causality (T8, T19). Network bursts occurring during workstation lock or idle intervals (>60s) trigger anomaly scoring.
- **Headless Server Profile:** Suspends HID input requirements. Causality shifts to parent-child service lineage, workload-specific request volume, and scheduled job baseline consistency.
- **Detector Telemetry Requirement:** Detectors must inspect host product type (`OSVERSIONINFOEX.wProductType != VER_NT_WORKSTATION` on Windows or headless display flags on Linux) to dynamically activate the appropriate baseline rule set.

8.6 Continuous-Domain Poisson Sampling vs. Timer Quantization

Discrete-interval sleep algorithms (`delay_ms = rand(100, 500)`) inevitably cluster into artificial time buckets, exposing timing generators to Power Spectral Density (PSD) Dirac combs and discrete timer quantization detectors (T3).

To eliminate discrete timer buckets, advanced traffic models evaluate continuous-domain inverse transform sampling:

$$T = -\ln(U) / \lambda$$

Where $U \sim \text{Uniform}(0, 1)$ is sampled as an IEEE 754 64-bit continuous floating-point value. Because continuous sampling spans an unquantized temporal continuum, discrete Fourier harmonics flatten into a continuous density envelope, depriving NDR sensors of discrete line artifacts. Defensive architectures must counter this not through scalar discretization, but through continuous statistical tail estimators and host kernel sleep inspection (T9, T18).

9. Maturity Model and Implementation Roadmap

9.1 Maturity Levels

Level	Definition	Characteristic Capabilities
L1 — Reactive	IOC and signature matching	Static domain/IP blocklists

Level	Definition	Characteristic Capabilities
L2 — Threshold	Volumetric and periodic alerts	Fixed rate limits, high noise
L3 — Statistical	Scaling, tail, and burst analytics	T1, T2, T5
L4 — Causal	Flow-to-process-to-identity binding	T6, T7, T8, T12
L5 — Adaptive & Hardware-Aware	Continuous calibration & hardware boundary auditing	T13, T15, T16–T20, automated baselining

9.2 Implementation Roadmap

- **0–30 Days (Foundational):** Deploy NetFlow/IPFIX scaling estimators (T1) and define metric baselines.
- **30–60 Days (Correlation):** Bind flow telemetry to host process lineage using PID and ProcessStartTime (T6, T12); deploy session validation (T7).
- **60–90 Days (Validation):** Establish defensive profile simulation (T15) and pilot deception flows (T13).
- **90–365 Days (Institutionalization):** Deploy long-horizon seasonality models (T14) and institutionalize quarterly calibration cycles.

10. Defensive Debt and Continuous Calibration

Detection engineering efficacy decays over time as adversaries adapt. Maintaining operational parity requires continuous multi-layer telemetry validation, quarterly baseline refreshes, and proactive telemetry dependency audits.

11. Conclusion

This report demonstrates that classical flow-level network statistics collapse against baseline-adapted timing evasion. Enterprise defense must execute three strategic imperatives:

1. Enrich flow records with sliding-window scaling and tail estimators.
2. Bind network flows causally to host kernel execution telemetry via collision-resistant composite keys.
3. Validate detection pipelines continuously using controlled simulation harnesses.

The detection question has evolved: *"Does this network traffic appear regular?"* is superseded by *"Is the host process and user session context generating this network flow consistent with the claimed identity?"*

12. Appendices

Appendix A — Mathematical Derivations

A.1 Point Event Train vs. IAT Sequence:

$$x(t) = \sum_n \delta(t - nT) \Rightarrow X(f) = (1/T) \cdot \sum_k \delta(f - k/T) \\ \text{IAT}_n = T \Rightarrow \text{DFT}\{X_{\text{IAT}}\}[0] = NT, \quad \text{DFT}\{X_{\text{IAT}}\}[k] = 0 \quad (k \neq 0)$$

A.2 Model A (Event-Time Jitter):

$$t_n = nT + \epsilon_n, \quad \epsilon_n \text{ i.i.d.}, \quad E[\epsilon] = 0, \quad \text{Var}(\epsilon) = \sigma_\epsilon^2 \\ \text{IAT}_n = T + \epsilon_n - \epsilon_{(n-1)} \\ \text{Var}(\text{IAT}_n) = 2\sigma_\epsilon^2 \\ \text{Cov}(\text{IAT}_n, \text{IAT}_{(n+1)}) = -\sigma_\epsilon^2 \Rightarrow \rho_1 = -0.5$$

A.3 Model B (Interval Jitter) and Renewal PSD:

$$\text{IAT}_n = T + U_n, \quad t_n = \sum \text{IAT}_i \\ S(f) = \lambda \cdot (1 - |\psi(f)|^2) / |1 - \psi(f)|^2 \quad (f \neq 0)$$

A.4 Jitter Model Summary:

Attribute	Model A (Event-Time Jitter)	Model B (Interval Jitter)		
Discrete Harmonics (k/T)	Present; scaled by λ	$\Phi_{\varepsilon}(k/T)$	2^k	Absent ($f \neq 0$); strictly continuous
Continuous Spectrum	Present; scaled by $\lambda(1 - \dots)$	$\Phi_{\varepsilon}(f)$	2^f	Present; defined by renewal PSD
Lag-1 Autocorrelation	$\rho_1 = -0.5$	$\rho_1 = 0$ (Independent)		

A.5 Exponential Surplus Variable:

$$\begin{aligned}
 X &= T + Y, \quad Y \sim \text{Exp}(\lambda) \\
 P(Y > y + t \mid Y > y) &= e^{(-\lambda t)} \\
 P(X > s + t \mid X > s) &= e^{(-\lambda t)}, \quad s \geq T
 \end{aligned}$$

A.6 Scaling and Heavy-Tailed Variance:

$$\begin{aligned}
 \text{Var}(X^m) &= \sigma^2 \cdot m^{(2H-2)} \\
 \text{ON/OFF Pareto Sources: } H &= (3 - \alpha) / 2, \quad 1 < \alpha < 2
 \end{aligned}$$

A.7 Hill Tail-Index Estimator:

$$\hat{\alpha}_{\text{Hill}}(k) = \left[\frac{1}{k} \cdot \sum_{i=1}^k \ln(x_{(i)} / x_{(k+1)}) \right]^{-1}$$

Appendix B — Glossary

Term	Operational Definition
IAT	Inter-Arrival Time between consecutive packets or flows
Event Train	Point process modeling exact occurrence timestamps
SRD / LRD	Short-Range / Long-Range Dependence in time series
MMPP	Markov-Modulated Poisson Process
Renewal Process	Point process with independent, identically distributed intervals
Palm–Khinchine	Theorem defining asymptotic convergence of superpositions to Poisson
fBm / fGn	Fractional Brownian Motion / Fractional Gaussian Noise
Hurst Exponent (H)	Self-similarity parameter ($0.5 < H < 1.0$ indicates LRD)
Tail Index (α)	Pareto shape parameter governing tail decay thickness
sinc	Unnormalized sinc function: $\sin(x) / x$
CCDF	Complementary Cumulative Distribution Function
ETW / eBPF	Windows Event Tracing / Linux extended Berkeley Packet Filters

Appendix C — MITRE ATT&CK Mapping Summary

Tactical ID	Primary ATT&CK Technique	Telemetry Scope
T1, T5, T14	T1071, T1029	Flow volume and scaling variance
T2, T11	T1030	Data transfer size constraints
T3	T1497.003	OS timer quantum artifacts
T4	N/A	Fleet-wide parameter clustering
T6, T12	N/A	Flow-to-process causal binding architecture
T7	T1053.005, T1543.003	Headless service execution verification
T8	N/A	Human input event correlation
T9	T1497.003	Thread sleep API invocation profiling
T10	T1071.001, T1573.002	TLS negotiation stack validation
T13, T15	N/A	Deception lures and calibration simulation

Tactical ID	Primary ATT&CK Technique	Telemetry Scope
T16	T1071, T1001.002	Storage subsystem pass-through auditing
T17	T1001	IOMMU and DMA coherent buffer verification
T18	T1497	PMU hardware memory barrier frequency
T19	T1027, T1071	Human-presence causality and idle correlation
T20	T1573.002, T1071.001	Subsystem TLS stack and ETW WinHTTP divergence

Appendix D — References

- Leland, W., Taqqu, M., Willinger, W., Wilson, D. (1994). *On the Self-Similar Nature of Ethernet Traffic*. IEEE/ACM Transactions on Networking.
- Paxson, V., Floyd, S. (1995). *Wide-Area Traffic: The Failure of Poisson Modeling*. IEEE/ACM Transactions on Networking.
- Willinger, W., Taqqu, M., Sherman, R., Wilson, D. (1997). *Self-Similarity Through High-Variability: Statistical Analysis of Ethernet LAN Traffic at the Source Level*. IEEE/ACM Transactions on Networking.
- Crovella, M., Bestavros, A. (1997). *Self-Similarity in World Wide Web Traffic: Evidence and Possible Causes*. IEEE/ACM Transactions on Networking.
- Fischer, W., Meier-Hellstern, K. (1993). *The Markov-Modulated Poisson Process (MMPP) Cookbook*. Performance Evaluation.
- Andersen, A., Nielsen, B. (1998). *A Markovian Approach for Modeling Packet Traffic with Long-Range Dependence*. IEEE JSAC.
- Norros, I. (1995). *On the Use of Fractional Brownian Motion in the Theory of Connectionless Networks*. IEEE JSAC.
- Taqqu, M., Teverovsky, V., Willinger, W. (1995). *Estimators for Long-Range Dependence: An Empirical Study*. Fractals.
- Hill, B. (1975). *A Simple General Approach to Inference About the Tail of a Distribution*. Annals of Statistics.
- Palm, C. (1943). *Intensitätsschwankungen im Fernsprechverkehr*. Ericsson Technics.
- Khinchin, A. (1955). *Mathematical Methods in the Theory of Queueing*. Griffin.
- Daley, D., Vere-Jones, D. (2003). *An Introduction to the Theory of Point Processes*. Springer.
- Feldmann, A., Gilbert, A., Willinger, W. (1998). *Data Networks as Cascades: Investigating the Multifractal Nature of Internet WAN Traffic*. ACM SIGCOMM.
- Abry, P., Veitch, D. (1998). *Wavelet Analysis of Long-Range-Dependent Traffic*. IEEE Transactions on Information Theory.
- Riedi, R., Crouse, M., Ribeiro, V., Baraniuk, R. (1999). *A Multifractal Wavelet Model with Application to Network Traffic*. IEEE Transactions on Information Theory.
- Kantelhardt, J., Zschiegner, S., Koscielny-Bunde, E., Havlin, S., Bunde, A., Stanley, H. (2002). *Multifractal Detrended Fluctuation Analysis of Nonstationary Time Series*. Physica A.
- Granger, C., Joyeux, R. (1980). *An Introduction to Long-Memory Time Series Models and Fractional Differencing*. Journal of Time Series Analysis.
- Hosking, J. (1981). *Fractional Differencing*. Biometrika.
- Baillie, R. (1996). *Long Memory Processes and Fractional Integration in Econometrics*. Journal of Econometrics.
- Krämer, W., Sibbertsen, P., Kleiber, C. (2002). *Long-Memory versus Structural Breaks: An Overview*. Statistical Papers.
- Künsch, H. (1989). *The Jackknife and the Bootstrap for General Stationary Observations*. Annals of Statistics.
- Politis, D., Romano, J. (1994). *The Stationary Bootstrap*. Journal of the American Statistical Association.
- Mirsky, Y., Doitshman, T., Elovici, Y., Shabtai, A. (2018). *Kitsune: An Ensemble of Autoencoders for Online Network Intrusion Detection*. NDSS.
- Boutaba, R., Salahuddin, M., Limam, N., et al. (2018). *A Comprehensive Survey on Machine Learning for Networking*. Journal of Internet Services and Applications.

- Lopez-Martin, M., Carro, B., et al. (2017). *Network Traffic Classifier With Convolutional and Recurrent Neural Networks for Internet of Things*. IEEE Access.
- Chalapathy, R., Chawla, S. (2019). *Deep Learning for Anomaly Detection: A Survey*. arXiv:1901.03407.
- MITRE Corporation. *ATT&CK Enterprise Matrix*.
- Microsoft. *Event Tracing for Windows: Kernel and Threat Intelligence Provider Documentation*.
- Linux Kernel Documentation. *BPF CO-RE and Tracepoint Programming Reference*.

Appendix E — Legal Notice & Intellectual Property Armor

E.1 Copyright and Ownership Declaration

This technical report, architectural analyses, mathematical frameworks, detection heuristics, source code, and empirical datasets are Copyright © 2026 Ogulcan (ByGhost) Erarslan. All rights reserved. Content is protected under international copyright treaties, the Berne Convention, and applicable intellectual property jurisprudence.

E.2 Licensing Scope and Constraints (CC BY-NC-ND 4.0)

This work is released strictly under the Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0) license:

- **Attribution:** Any academic citation or non-commercial reference must preserve explicit credit to the author (Ogulcan 'ByGhost' Erarslan), document title, identifier (SOC-NDR-MASTER-001-EN), and verified repositories (<https://byghost.tr/> and <https://github.com/ByGh00st>).
- **Non-Commercial:** The technical report, detection mechanics, mathematical proofs, and architectural models cannot be utilized for commercial purposes, revenue-generating operations, or monetized advisory packages.
- **No Derivatives:** The material cannot be modified, remixed, transformed, decomposed, or built upon for public distribution without separate bilateral licensing.

E.3 Absolute Prohibition on Commercial & EDR/NDR Integration

Without an explicit, verifiable written commercial licensing agreement executed by the copyright holder:

1. Direct or indirect telemetry extraction, porting of tactical algorithms (T1–T20), eBPF kernels, or detection heuristics into commercial cybersecurity platforms (EDR, NDR, SIEM, SOAR, XDR) is strictly prohibited.
2. Providing academic attribution does NOT confer commercial distribution or software embedding rights.
3. Stripping metadata, cover pages, or running footers to repackage this report under corporate white-label branding constitutes willful copyright infringement and theft of trade assets.

E.4 Jurisdiction and Enforcement

Any breach of these terms shall be governed by and construed in accordance with international intellectual property treaties, with exclusive jurisdiction resting in the specialized Intellectual and Industrial Property Rights Courts. The author reserves all rights to seek preliminary injunctions, statutory damages, and restitution of illicit commercial revenues.

BYGHOST SECURITY

Oğulcan (ByGhost) Erarslan

Senior Solutions Architect | Cyber Security Specialist

ByGhost Security · <https://byghost.tr/> · github.com/ByGh00st · [linkedin.com/in/byghost-tr](https://www.linkedin.com/in/byghost-tr)